



Code Snippet Classification

Big Data Computing Project

Andrea Trianni - Emanuele Mercanti

The Task



Given a **short** snippet of code, predict the programming language.



Useful for Visual IDE tool :

- to highlight keywords
- to check syntax error

Useful for Online forums :

- clusters un-tagged questions

The environment



Development



Training stage

The Dataset





GitHub SQL Dataset

Version available: FULL (60GB) and LITE (3GB)


20 Different languages covered + Unknown Class

Length of snippets is 5 rows

We have access also to other Fields, like License Type


Visualize the data

class sizes




Visualize the data

snippet length distribution



Visualize the data

boxenplot: snippet length distribution by class



Our Approach

The preprocessing pipeline





Cleaning the data

Manual **stratify split** in train, dev and test set

Remove **non-programming** languages (csv, json, ..)

Remove **noisy unknown** class

Merge **Shell** family languages (shell, bash, ..)

Remove **outlier** rows

Check for **missing** data

Calculating class **weights**



Tokenize the snippet

Our first try:

A small Vocabulary manually built from a fixed set of keywords + punctuation

Tokenizing finding substring matches, using the vocabulary.

Strength:

Low dimensionality of vectors helps simple models to achieve high performance

Weakness:

Powerful models not able to reach their best performances

No parallelization possible with simple implementation



Tokenize the snippet

So, we rely on Spark **RegexTokenizer** :

```
( [\\n\\t] | ([A-Za-z_]+\\b) | [!\\#\\$%\\&\\*\\+\\:~-\\. /<=>\\?@\\^_\\|\\~]+ | [ \\(\\) , ; \\{\\}\\[\\] `\"' ] )
```



Indentation



Keywords + Var.



Operators



Brackets



Encoding the snippet

We tried 2 encoding strategy :



BOW



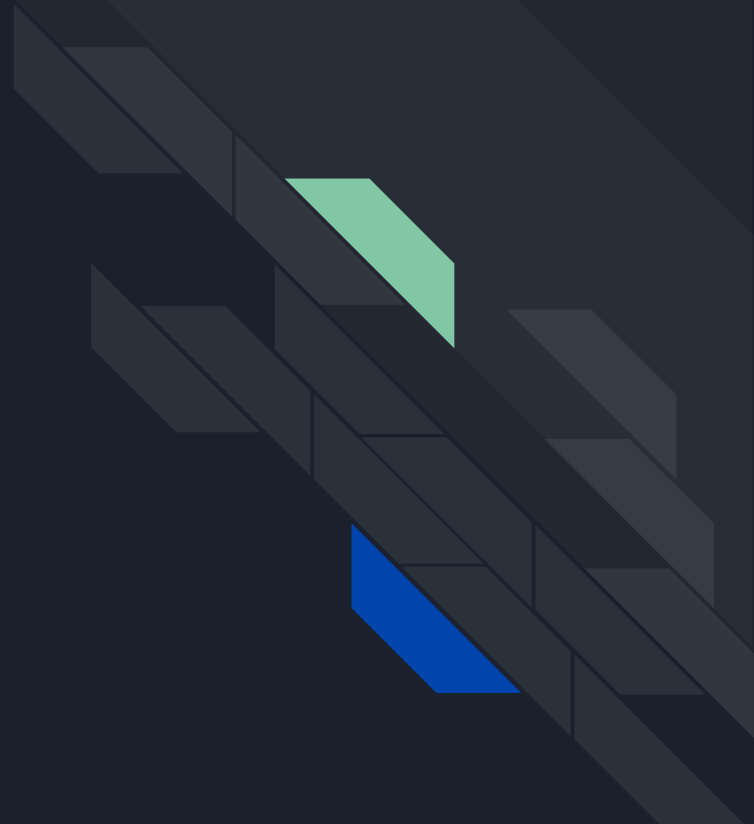
TF-IDF

5k Max Vocabulary Size

Sometimes, less is more.

Our Approach

The training stage





Training the models

We train and evaluate the following models:

Decision tree classifier

Gradient-boosted tree classifier [OVR] (failed)

Random forest

Multinomial logistic regression

Naive Bayes classifier

Linear Support vectors machines [OVR]



Evaluating the models

We obtain the following results:

MODEL	F1 Score (dev)	Accuracy (dev)
Decision tree classifier	0.38	0.45
Gradient-boosted tree	?	?
Random forest	0.55	0.53
Multinomial logistic reg.	0.83	0.83
Naive bayes	0.74	0.73
Linear SVM	0.82	0.82



Results explanation

Winners

Logistic regression Linear SVM



Data is linearly separable

Data can be modeled by a multinomial distribution

Losers

Decision tree Random Forest



Data is very sparse

Data might not be aligned to axis



Benchmark the model

So, we select the best model and we benchmark it on **test** set, using:

Cross Validation + Grid Search

Accuracy + F1 score

Confusion matrix

Learning curve

Results



Benchmarks:

(Grid Search) 5 Fold Cross-Validation F1: **82.5%**

Test set results:


F1-Score: **82.2 %**

Accuracy: **82.0%**

Confusion Matrix



Learning Curve



Web Application





How Does it Works?

Instead of leveraging Docker (local) and build a Web Application using Flask REST API..

We decided to use Colab Runtime as a virtual machine, configure Spark with a Master Node
and deploy our application using **Streamlit**

Actually, you can try to use it (the link will be provided during the live presentation)

Code snippets classification

Enter here your snippet..

Or choose a file



Drag and drop file here

Limit 200MB per file

Browse files



sample.cpp 86.0B



Submit



Thank you!